

EDA Data Transformation

November 27, 2024

1 Data Transformation

1. Data deduplication involves the identification of duplicates and their removal.
2. Key restructuring involves transforming any keys with built-in meanings to the generic keys.
3. Data cleansing involves extracting words and deleting out-of-date, inaccurate, and incomplete information from the source language without extracting the meaning or information to enhance the accuracy of the source data.
4. Data validation is a process of formulating rules or algorithms that help in validating different types of data against some known issues.
5. Format revisioning involves converting from one format to another.
6. Data derivation consists of creating a set of rules to generate more information from the data source.
7. Data aggregation involves searching, extracting, summarizing, and preserving important information in different types of reporting systems.
8. Data integration involves converting different data types and merging them into a common structure or schema.
9. Data filtering involves identifying information relevant to any particular user.
10. Data joining involves establishing a relationship between two or more tables.

```
[1]: import pandas as pd
import numpy as np
```

2 Combining dataframes

```
[2]: dataframe1 = pd.DataFrame({'StudentID': [1, 3, 5, 7, 9, 11, 13, 15, 17, 19,
↪21, 23, 25, 27, 29], 'Score' : [89, 39, 50, 97, 22, 66, 31, 51, 71, 91, 56,
↪32, 52, 73, 92]})
dataframe2 = pd.DataFrame({'StudentID': [2, 4, 6, 8, 10, 12, 14, 16, 18, 20,
↪22, 24, 26, 28, 30], 'Score': [98, 93, 44, 77, 69, 56, 31, 53, 78, 93, 56,
↪77, 33, 56, 27]})

# In the dataset above, the first column contains information about student_
↪identifier and the second column contains their respective scores in any_
↪subject. The structure of the dataframes is same in the bothe case. In this_
↪case, we would need to concatenate both of them.
```

```
[3]: # We can do that by using Pandas concat() method.
```

```
dataframe = pd.concat([dataFrame1, dataFrame2], ignore_index=True)
dataframe
```

```
[3]:
```

	StudentID	Score
0	1	89
1	3	39
2	5	50
3	7	97
4	9	22
5	11	66
6	13	31
7	15	51
8	17	71
9	19	91
10	21	56
11	23	32
12	25	52
13	27	73
14	29	92
15	2	98
16	4	93
17	6	44
18	8	77
19	10	69
20	12	56
21	14	31
22	16	53
23	18	78
24	20	93
25	22	56
26	24	77
27	26	33
28	28	56
29	30	27

The argument `ignore_index` creates new index and its absence keeps the original indices. Note, we combined the dataframes along `axis=0`, that is to say, we combined together along same direction. What if we want to combine both side by side. Then we have to specify `axis = 1`. Check the output and see the difference.

```
[4]: pd.concat([dataFrame1, dataFrame2], axis=1)
```

```
[4]:
```

	StudentID	Score	StudentID	Score
0	1	89	2	98
1	3	39	4	93
2	5	50	6	44

3	7	97	8	77
4	9	22	10	69
5	11	66	12	56
6	13	31	14	31
7	15	51	16	53
8	17	71	18	78
9	19	91	20	93
10	21	56	22	56
11	23	32	24	77
12	25	52	26	33
13	27	73	28	56
14	29	92	30	27

3 Merging

In the first example, you received two files for same subject. Now, consider the use case where you are teaching two courses. So, you will get two dataframes from each sections: two for Software engineering course and another two for Introduction to Machine learning course. Check the figure given below:

```
[5]: df1SE = pd.DataFrame({ 'StudentID': [9, 11, 13, 15, 17, 19, 21, 23, 25, 27, 29], 'ScoreSE' : [22, 66, 31, 51, 71, 91, 56, 32, 52, 73, 92]})
df2SE = pd.DataFrame({'StudentID': [2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30], 'ScoreSE': [98, 93, 44, 77, 69, 56, 31, 53, 78, 93, 56, 77, 33, 56, 27]})

df1ML = pd.DataFrame({ 'StudentID': [1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21, 23, 25, 27, 29], 'ScoreML' : [39, 49, 55, 77, 52, 86, 41, 77, 73, 51, 86, 82, 92, 23, 49]})
df2ML = pd.DataFrame({'StudentID': [2, 4, 6, 8, 10, 12, 14, 16, 18, 20], 'ScoreML': [93, 44, 78, 97, 87, 89, 39, 43, 88, 78]})
```

As you can see in the dataset above, you have two dataframes for each subjects. So the first task would be to concatenate these two subjects into one. Secondly, these students have taken Introduction to Machine Learning course as well. So, we need to merge these score into the same dataframes. There are several ways to do this. Let us explore some options.

```
[6]: # Option 1
dfSE = pd.concat([df1SE, df2SE], ignore_index=True)
dfML = pd.concat([df1ML, df2ML], ignore_index=True)

df = pd.concat([dfML, dfSE], axis=1)
df
```

```
[6]:
```

	StudentID	ScoreML	StudentID	ScoreSE
0	1.0	39.0	9	22
1	3.0	49.0	11	66

2	5.0	55.0	13	31
3	7.0	77.0	15	51
4	9.0	52.0	17	71
5	11.0	86.0	19	91
6	13.0	41.0	21	56
7	15.0	77.0	23	32
8	17.0	73.0	25	52
9	19.0	51.0	27	73
10	21.0	86.0	29	92
11	23.0	82.0	2	98
12	25.0	92.0	4	93
13	27.0	23.0	6	44
14	29.0	49.0	8	77
15	2.0	93.0	10	69
16	4.0	44.0	12	56
17	6.0	78.0	14	31
18	8.0	97.0	16	53
19	10.0	87.0	18	78
20	12.0	89.0	20	93
21	14.0	39.0	22	56
22	16.0	43.0	24	77
23	18.0	88.0	26	33
24	20.0	78.0	28	56
25	NaN	NaN	30	27

[7]: *# Option 2*

```
dfSE = pd.concat([df1SE, df2SE], ignore_index=True)
```

```
dfML = pd.concat([df1ML, df2ML], ignore_index=True)
```

```
df = dfSE.merge(dfML, how='inner')
```

```
df
```

Here, you will perform inner join with each dataframe. That is to say, if an
→ item exists on the both dataframe, will be included in the new dataframe.
→ This means, we will get the list of students who are appearing in both the
→ courses.

[7]:

	StudentID	ScoreSE	ScoreML
--	-----------	---------	---------

0	9	22	52
1	11	66	86
2	13	31	41
3	15	51	77
4	17	71	73
5	19	91	51
6	21	56	86
7	23	32	82
8	25	52	92

9	27	73	23
10	29	92	49
11	2	98	93
12	4	93	44
13	6	44	78
14	8	77	97
15	10	69	87
16	12	56	89
17	14	31	39
18	16	53	43
19	18	78	88
20	20	93	78

```
[8]: # Option 3
dfSE = pd.concat([df1SE, df2SE], ignore_index=True)
dfML = pd.concat([df1ML, df2ML], ignore_index=True)

df = dfSE.merge(dfML, how='left')
df
```

```
[8]:
```

	StudentID	ScoreSE	ScoreML
0	9	22	52.0
1	11	66	86.0
2	13	31	41.0
3	15	51	77.0
4	17	71	73.0
5	19	91	51.0
6	21	56	86.0
7	23	32	82.0
8	25	52	92.0
9	27	73	23.0
10	29	92	49.0
11	2	98	93.0
12	4	93	44.0
13	6	44	78.0
14	8	77	97.0
15	10	69	87.0
16	12	56	89.0
17	14	31	39.0
18	16	53	43.0
19	18	78	88.0
20	20	93	78.0
21	22	56	NaN
22	24	77	NaN
23	26	33	NaN
24	28	56	NaN
25	30	27	NaN

```
[9]: # Option 4
dfSE = pd.concat([df1SE, df2SE], ignore_index=True)
dfML = pd.concat([df1ML, df2ML], ignore_index=True)

df = dfSE.merge(dfML, how='right')
df
```

```
[9]:
```

	StudentID	ScoreSE	ScoreML
0	1	NaN	39
1	3	NaN	49
2	5	NaN	55
3	7	NaN	77
4	9	22.0	52
5	11	66.0	86
6	13	31.0	41
7	15	51.0	77
8	17	71.0	73
9	19	91.0	51
10	21	56.0	86
11	23	32.0	82
12	25	52.0	92
13	27	73.0	23
14	29	92.0	49
15	2	98.0	93
16	4	93.0	44
17	6	44.0	78
18	8	77.0	97
19	10	69.0	87
20	12	56.0	89
21	14	31.0	39
22	16	53.0	43
23	18	78.0	88
24	20	93.0	78

```
[10]: # Option 5
dfSE = pd.concat([df1SE, df2SE], ignore_index=True)
dfML = pd.concat([df1ML, df2ML], ignore_index=True)

df = dfSE.merge(dfML, how='outer')
df
```

```
[10]:
```

	StudentID	ScoreSE	ScoreML
0	1	NaN	39.0
1	2	98.0	93.0
2	3	NaN	49.0
3	4	93.0	44.0
4	5	NaN	55.0

5	6	44.0	78.0
6	7	NaN	77.0
7	8	77.0	97.0
8	9	22.0	52.0
9	10	69.0	87.0
10	11	66.0	86.0
11	12	56.0	89.0
12	13	31.0	41.0
13	14	31.0	39.0
14	15	51.0	77.0
15	16	53.0	43.0
16	17	71.0	73.0
17	18	78.0	88.0
18	19	91.0	51.0
19	20	93.0	78.0
20	21	56.0	86.0
21	22	56.0	NaN
22	23	32.0	82.0
23	24	77.0	NaN
24	25	52.0	92.0
25	26	33.0	NaN
26	27	73.0	23.0
27	28	56.0	NaN
28	29	92.0	49.0
29	30	27.0	NaN

```
[88]: #Merging on index
left1 = pd.DataFrame({'key': ['apple','ball','apple', 'apple','ball', 'cat'],
↳ 'value': range(6)})
right1 = pd.DataFrame({'group_val': [33.4, 5]}, index=['apple', 'ball'])
left1
```

```
[88]:      key  value
0  apple      0
1   ball      1
2  apple      2
3  apple      3
4   ball      4
5    cat      5
```

```
[87]: right1
```

```
[87]:      group_val
apple      33.4
ball       5.0
```

```
[89]: #let's try merging using an inner join, which is the default type of merge.
df = pd.merge(left1, right1, left_on='key', right_index=True)
df
```

```
[89]:      key  value  group_val
0  apple      0      33.4
1   ball      1       5.0
2  apple      2      33.4
3  apple      3      33.4
4   ball      4       5.0
```

```
[90]: # let's try merging using an outer join, as follows
df = pd.merge(left1, right1, left_on='key', right_index=True, how='outer')
df
```

```
[90]:      key  value  group_val
0  apple      0      33.4
2  apple      2      33.4
3  apple      3      33.4
1   ball      1       5.0
4   ball      4       5.0
5    cat      5       NaN
```

3.1 Reshaping and pivoting

```
[91]: data = np.arange(15).reshape((3,5))
indexers = ['Rainfall', 'Humidity', 'Wind']
dframe1 = pd.DataFrame(data, index=indexers, columns=['Bergen', 'Oslo', 'Trondheim', 'Stavanger', 'Kristiansand'])
dframe1
```

```
[91]:      Bergen  Oslo  Trondheim  Stavanger  Kristiansand
Rainfall      0     1           2           3           4
Humidity      5     6           7           8           9
Wind          10    11          12          13          14
```

```
[92]: # using the stack() method on the preceding dframe1, we can pivot the columns
      ↪ into rows to produce a series
stacked = dframe1.stack()
stacked
```

```
[92]: Rainfall  Bergen      0
      Oslo      1
      Trondheim  2
      Stavanger  3
      Kristiansand 4
Humidity  Bergen      5
```



```

        Oslo        6
        Trondheim   7
        Stavanger   8
        Kristiansand 9
Wind      Bergen    10
        Oslo        11
        Trondheim   12
        Stavanger   13
        Kristiansand 14
dtype: int64

```

```

[93]: # The preceding series stored unstacked in the variable can be rearranged into
      ↪ a dataframe using the unstack() method:
stacked.unstack()

```

```

[93]:      Bergen  Oslo  Trondheim  Stavanger  Kristiansand
Rainfall      0     1           2           3           4
Humidity      5     6           7           8           9
Wind          10    11          12          13          14

```

```

[94]: # Note that there is a chance that unstacking will create missing data if all
      ↪ the values are not present in each of the sub-groups. ex:
series1 = pd.Series([000, 111, 222, 333], index=['zeros', 'ones', 'twos',
      ↪ 'threes'])
series2 = pd.Series([444, 555, 666], index=['fours', 'fives', 'sixes'])
frame2 = pd.concat([series1, series2], keys=['Number1', 'Number2'])

frame2.unstack()

```

```

[94]:      fives  fours   ones  sixes  threes  twos  zeros
Number1   NaN   NaN  111.0   NaN   333.0  222.0   0.0
Number2  555.0  444.0   NaN  666.0   NaN   NaN   NaN

```

3.2 Transformatio techniques

3.2.1 Performing data deduplication

```

[95]: frame3 = pd.DataFrame({'column 1': ['Looping'] * 3 + ['Functions'] * 4, 'column_
      ↪ 2': [10, 10, 22, 23, 23, 24, 24]})
frame3

```

```

[95]:      column 1  column 2
0      Looping      10
1      Looping      10
2      Looping      22
3  Functions      23
4  Functions      23

```

```
5 Functions      24
6 Functions      24
```

```
[96]: frame3.duplicated()
```

```
[96]: 0    False
      1     True
      2    False
      3    False
      4     True
      5    False
      6     True
      dtype: bool
```

```
[97]: frame4 = frame3.drop_duplicates()
      frame4
```

```
[97]:   column 1  column 2
0    Looping      10
2    Looping      22
3  Functions      23
5  Functions      24
```

```
[98]: frame3['column 3'] = range(7)
      frame5 = frame3.drop_duplicates(['column 2'])
      frame5
```

```
[98]:   column 1  column 2  column 3
0    Looping      10         0
2    Looping      22         2
3  Functions      23         3
5  Functions      24         5
```

3.3 Replacing values

```
[100]: import numpy as np
      replaceFrame = pd.DataFrame({'column 1': [200., 3000., -786., 3000., 234., 444.,
      ↪, -786., 332., 3332. ], 'column 2': range(9)})
      replaceFrame.replace(to_replace = -786, value= np.nan)
```

```
[100]:   column 1  column 2
0     200.0         0
1    3000.0         1
2         NaN         2
3    3000.0         3
4     234.0         4
5     444.0         5
```

6	NaN	6
7	332.0	7
8	3332.0	8

```
[101]: replaceFrame = pd.DataFrame({'column 1': [200., 3000., -786., 3000., 234., 444.,
↪, -786., 332., 3332. ], 'column 2': range(9)})
replaceFrame.replace(to_replace =[-786, 0], value= [np.nan, 2])
```

```
[101]:
```

	column 1	column 2
0	200.0	2
1	3000.0	1
2	NaN	2
3	3000.0	3
4	234.0	4
5	444.0	5
6	NaN	6
7	332.0	7
8	3332.0	8

3.4 Handling missing data

Whenever there are missing values, a NaN value is used, which indicates that there is no value specified for that particular index. There could be several reasons why a value could be NaN: 1. It can happen when data is retrieved from an external source and there are some incomplete values in the dataset. 2. It can also happen when we join two different datasets and some values are not matched. 3. Missing values due to data collection errors. 4. When the shape of data changes, there are new additional rows or columns that are not determined. 5. Reindexing of data can result in incomplete data.

```
[102]: data = np.arange(15, 30).reshape(5, 3)
dfx = pd.DataFrame(data, index=['apple', 'banana', 'kiwi', 'grapes', 'mango'],
↪columns=['store1', 'store2', 'store3'])
dfx
```

```
[102]:
```

	store1	store2	store3
apple	15	16	17
banana	18	19	20
kiwi	21	22	23
grapes	24	25	26
mango	27	28	29

```
[103]: dfx['store4'] = np.nan
dfx.loc['watermelon'] = np.arange(15, 19)
dfx.loc['oranges'] = np.nan
dfx['store5'] = np.nan
dfx['store4']['apple'] = 20.
dfx
```

/tmp/ipykernel_44487/320494991.py:5: FutureWarning: ChainedAssignmentError: behaviour will change in pandas 3.0!

You are setting values through chained assignment. Currently this works in certain cases, but when using Copy-on-Write (which will become the default behaviour in pandas 3.0) this will never work to update the original DataFrame or Series, because the intermediate object on which we are setting values will behave as a copy.

A typical example is when you are setting values in a column of a DataFrame, like:

```
df["col"][row_indexer] = value
```

Use `df.loc[row_indexer, "col"] = values` instead, to perform the assignment in a single step and ensure this keeps updating the original `df`.

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
dfx['store4']['apple'] = 20.
```

```
[103]:
```

	store1	store2	store3	store4	store5
apple	15.0	16.0	17.0	20.0	NaN
banana	18.0	19.0	20.0	NaN	NaN
kiwi	21.0	22.0	23.0	NaN	NaN
grapes	24.0	25.0	26.0	NaN	NaN
mango	27.0	28.0	29.0	NaN	NaN
watermelon	15.0	16.0	17.0	18.0	NaN
oranges	NaN	NaN	NaN	NaN	NaN

```
[ ]:
```

```
[105]: dfx.isnull()
```

```
[105]:
```

	store1	store2	store3	store4	store5
apple	False	False	False	False	True
banana	False	False	False	True	True
kiwi	False	False	False	True	True
grapes	False	False	False	True	True
mango	False	False	False	True	True
watermelon	False	False	False	False	True
oranges	True	True	True	True	True

```
[106]: # We can use the sum() method to count the number of NaN values in each store.  
dfx.isnull().sum()
```

```
[106]: store1    1  
store2    1
```

```
store3    1
store4    5
store5    7
dtype: int64
```

```
[107]: dfx.isnull().sum().sum()
```

```
[107]: 15
```

```
[109]: # instead of counting the number of missing values, we can count the number of
      ↪ reported values:
      dfx.count()
```

```
[109]: store1    6
      store2    6
      store3    6
      store4    2
      store5    0
      dtype: int64
```

3.5 Dropping missing values

```
[111]: #determine the null values
      dfx.store4[dfx.store4.notnull()]
```

```
[111]: apple          20.0
      watermelon     18.0
      Name: store4, dtype: float64
```

```
[112]: # Now, we can use the dropna() method to remove the rows:
      dfx.store4.dropna()
```

```
[112]: apple          20.0
      watermelon     18.0
      Name: store4, dtype: float64
```

```
[116]: dfx
```

```
[116]:
```

	store1	store2	store3	store4	store5
apple	15.0	16.0	17.0	20.0	NaN
banana	18.0	19.0	20.0	NaN	NaN
kiwi	21.0	22.0	23.0	NaN	NaN
grapes	24.0	25.0	26.0	NaN	NaN
mango	27.0	28.0	29.0	NaN	NaN
watermelon	15.0	16.0	17.0	18.0	NaN
oranges	NaN	NaN	NaN	NaN	NaN

```
[117]: # Note that the dropna() method just returns a copy of the dataframe by
        ↪dropping the rows with NaN. The original dataframe is not changed.
        # If dropna() is applied to the entire dataframe, then it will drop all the
        ↪rows from the dataframe, because there is at least one NaN value in our
        ↪dataframe:
        # dfx.dropna()

        # Dropping by rows
        dfx.dropna(how='all')
```

```
[117]:
```

	store1	store2	store3	store4	store5
apple	15.0	16.0	17.0	20.0	NaN
banana	18.0	19.0	20.0	NaN	NaN
kiwi	21.0	22.0	23.0	NaN	NaN
grapes	24.0	25.0	26.0	NaN	NaN
mango	27.0	28.0	29.0	NaN	NaN
watermelon	15.0	16.0	17.0	18.0	NaN

```
[118]: # Dropping by columns
        dfx.dropna(how='all', axis=1)
```

```
[118]:
```

	store1	store2	store3	store4
apple	15.0	16.0	17.0	20.0
banana	18.0	19.0	20.0	NaN
kiwi	21.0	22.0	23.0	NaN
grapes	24.0	25.0	26.0	NaN
mango	27.0	28.0	29.0	NaN
watermelon	15.0	16.0	17.0	18.0
oranges	NaN	NaN	NaN	NaN

3.5.1 Mathematical operations with NaN

Note the following things: 1. When a NumPy function encounters NaN values, it returns NaN. 2. Pandas, on the other hand, ignores the NaN values and moves ahead with processing. When performing the sum operation, NaN is treated as 0. If all the values are NaN, the result is also NaN.

```
[120]: ar1 = np.array([100, 200, np.nan, 300])
        ser1 = pd.Series(ar1)
        ser1
```

```
[120]: 0    100.0
        1    200.0
        2     NaN
        3    300.0
        dtype: float64
```

```
[122]: ar1.mean(), ser1.mean()
```

```
[122]: (nan, 200.0)
```

3.6 Filling missing values

```
[123]: filledDf = dfx.fillna(0)
filledDf
```

```
[123]:
```

	store1	store2	store3	store4	store5
apple	15.0	16.0	17.0	20.0	0.0
banana	18.0	19.0	20.0	0.0	0.0
kiwi	21.0	22.0	23.0	0.0	0.0
grapes	24.0	25.0	26.0	0.0	0.0
mango	27.0	28.0	29.0	0.0	0.0
watermelon	15.0	16.0	17.0	18.0	0.0
oranges	0.0	0.0	0.0	0.0	0.0

```
[124]: dfx.mean()
```

```
[124]: store1    20.0
store2    21.0
store3    22.0
store4    19.0
store5     NaN
dtype: float64
```

```
[125]: filledDf.mean()
```

```
[125]: store1    17.142857
store2    18.000000
store3    18.857143
store4     5.428571
store5     0.000000
dtype: float64
```

```
[128]: # Backward and forward filling
# Here, from the forward-filling technique, the last known value is 20 and
# hence the rest of the NaN values are replaced by it.
dfx.store4.fillna(method='ffill')
```

```
/tmp/ipykernel_44487/1459500098.py:3: FutureWarning: Series.fillna with 'method'
is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill()
instead.
```

```
dfx.store4.fillna(method='ffill')
```

```
[128]: apple      20.0
      banana     20.0
      kiwi       20.0
      grapes     20.0
      mango      20.0
      watermelon 18.0
      oranges    18.0
      Name: store4, dtype: float64
```

```
[129]: #The direction of the fill can be changed by changing method='bfill'. Check the
      ↪ following example: ser3 = pd.Series([100, np.nan, np.nan, np.nan, 292])
      ser3.interpolate()
      dfx.store4.fillna(method='bfill')
```

/tmp/ipykernel_44487/1648656597.py:2: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.

```
dfx.store4.fillna(method='bfill')
```

```
[129]: apple      20.0
      banana     18.0
      kiwi       18.0
      grapes     18.0
      mango      18.0
      watermelon 18.0
      oranges     NaN
      Name: store4, dtype: float64
```

```
[131]: # Interpolating missing values
      ser3 = pd.Series([100, np.nan, np.nan, np.nan, 292])
      ser3.interpolate()
```

```
[131]: 0    100.0
      1    148.0
      2    196.0
      3    244.0
      4    292.0
      dtype: float64
```

3.6.1 Renaming axis indexes

```
[132]: dframe1.index = dframe1.index.map(str.upper)
      dframe1
```

```
[132]:
```

	Bergen	Oslo	Trondheim	Stavanger	Kristiansand
RAINFALL	0	1	2	3	4
HUMIDITY	5	6	7	8	9

WIND	10	11	12	13	14
------	----	----	----	----	----

```
[133]: dframe1.rename(index=str.title, columns=str.upper)
```

```
[133]:
```

	BERGEN	OSLO	TRONDHEIM	STAVANGER	KRISTIANSAND
Rainfall	0	1	2	3	4
Humidity	5	6	7	8	9
Wind	10	11	12	13	14

3.7 Outlier detection and filtering

Outliers are data points that diverge from other observations for several reasons. During the EDA phase, one of our common tasks is to detect and filter these outliers. The main reason for this detection and filtering of outliers is that the presence of such outliers can cause serious issues in statistical analysis.

```
[11]: df = pd.read_csv('https://raw.githubusercontent.com/PacktPublishing/
↳hands-on-exploratory-data-analysis-with-python/master/Chapter%204/sales.csv')
df.head(10)
```

```
[11]:
```

	Account	Company	Order	SKU	Country	Year	\
0	123456779	Kulas Inc	99985	s9-supercomputer	Aruba	1981	
1	123456784	Gitbub	99986	s4-supercomputer	Brazil	2001	
2	123456782	Kulas Inc	99990	s10-supercomputer	Montserrat	1973	
3	123456783	My SQ Man	99999	s1-supercomputer	El Salvador	2015	
4	123456787	ABC Dogma	99996	s6-supercomputer	Poland	1970	
5	123456778	Super Sexy Dingo	99996	s9-supercomputer	Costa Rica	2004	
6	123456783	ABC Dogma	99981	s11-supercomputer	Spain	2006	
7	123456785	ABC Dogma	99998	s9-supercomputer	Belarus	2015	
8	123456778	Loolo INC	99997	s8-supercomputer	Mauritius	1999	
9	123456775	Kulas Inc	99997	s7-supercomputer	French Guiana	2004	

	Quantity	UnitPrice	transactionComplete
0	5148	545	False
1	3262	383	False
2	9119	407	True
3	3097	615	False
4	3356	91	True
5	2474	136	True
6	4081	195	False
7	6576	603	False
8	2460	36	False
9	1831	664	True

```
[12]: #@title Default title text
#Add new colum that is the total price based on the quantity and the unit price
```

```
df['TotalPrice'] = df['UnitPrice'] * df['Quantity']
df.head(10)
```

```
[12]:
```

	Account	Company	Order	SKU	Country	Year	\
0	123456779	Kulas Inc	99985	s9-supercomputer	Aruba	1981	
1	123456784	GitHub	99986	s4-supercomputer	Brazil	2001	
2	123456782	Kulas Inc	99990	s10-supercomputer	Montserrat	1973	
3	123456783	My SQ Man	99999	s1-supercomputer	El Salvador	2015	
4	123456787	ABC Dogma	99996	s6-supercomputer	Poland	1970	
5	123456778	Super Sexy Dingo	99996	s9-supercomputer	Costa Rica	2004	
6	123456783	ABC Dogma	99981	s11-supercomputer	Spain	2006	
7	123456785	ABC Dogma	99998	s9-supercomputer	Belarus	2015	
8	123456778	Loolo INC	99997	s8-supercomputer	Mauritius	1999	
9	123456775	Kulas Inc	99997	s7-supercomputer	French Guiana	2004	

	Quantity	UnitPrice	transactionComplete	TotalPrice
0	5148	545	False	2805660
1	3262	383	False	1249346
2	9119	407	True	3711433
3	3097	615	False	1904655
4	3356	91	True	305396
5	2474	136	True	336464
6	4081	195	False	795795
7	6576	603	False	3965328
8	2460	36	False	88560
9	1831	664	True	1215784

```
[13]: df['Company'].value_counts()
```

```
[13]:
```

Company	
My SQ Man	869
Kirlosker Service Center	863
Will LLC	862
ABC Dogma	848
Kulas Inc	840
Gen Power	836
Name IT	836
Super Sexy Dingo	828
GitHub	823
Loolo INC	822
SAS Web Tec	798
Pryanika Ji	775

Name: count, dtype: int64

```
[14]: df.describe()
```

```
[14]:
```

	Account	Order	Year	Quantity	UnitPrice \
count	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000
mean	1.234568e+08	99989.562900	1994.619800	4985.447300	355.866600
std	5.741156e+00	5.905551	14.432771	2868.949686	201.378478
min	1.234568e+08	99980.000000	1970.000000	0.000000	10.000000
25%	1.234568e+08	99985.000000	1982.000000	2505.750000	181.000000
50%	1.234568e+08	99990.000000	1995.000000	4994.000000	356.000000
75%	1.234568e+08	99995.000000	2007.000000	7451.500000	531.000000
max	1.234568e+08	99999.000000	2019.000000	9999.000000	700.000000

	TotalPrice
count	1.000000e+04
mean	1.773301e+06
std	1.540646e+06
min	0.000000e+00
25%	5.003370e+05
50%	1.335698e+06
75%	2.711653e+06
max	6.841580e+06

3.8 Reshaping with Hierarchical Indexing

```
[15]: data = np.arange(15).reshape((3,5))
indexers = ['Rainfall', 'Humidity', 'Wind']
dframe1 = pd.DataFrame(data, index=indexers, columns=['Bergen', 'Oslo', 'Trondheim', 'Stavanger', 'Kristiansand'])
dframe1
```

```
[15]:
```

	Bergen	Oslo	Trondheim	Stavanger	Kristiansand
Rainfall	0	1	2	3	4
Humidity	5	6	7	8	9
Wind	10	11	12	13	14

```
[16]: stacked = dframe1.stack()
stacked
```

```
[16]:
```

Rainfall	Bergen	0
	Oslo	1
	Trondheim	2
	Stavanger	3
	Kristiansand	4
Humidity	Bergen	5
	Oslo	6
	Trondheim	7
	Stavanger	8
Wind	Kristiansand	9
	Bergen	10

```

        Oslo          11
        Trondheim     12
        Stavanger     13
        Kristiansand  14
dtype: int64

```

```
[17]: stacked.unstack()
```

```
[17]:
```

	Bergen	Oslo	Trondheim	Stavanger	Kristiansand
Rainfall	0	1	2	3	4
Humidity	5	6	7	8	9
Wind	10	11	12	13	14

```
[18]: series1 = pd.Series([000, 111, 222, 333], index=['zeros', 'ones', 'twos', 'threes'])
      series2 = pd.Series([444, 555, 666], index=['fours', 'fives', 'sixs'])

      frame2 = pd.concat([series1, series2], keys=['Number1', 'Number2'])
      frame2.unstack()
```

```
[18]:
```

	fives	fours	ones	sixs	threes	twos	zeros
Number1	NaN	NaN	111.0	NaN	333.0	222.0	0.0
Number2	555.0	444.0	NaN	666.0	NaN	NaN	NaN

3.9 Forward and backward filling of the missing values

```
[49]: dfx.store4.fillna(method='ffill')
```

```

/tmp/ipykernel_44487/4057730406.py:1: FutureWarning: Series.fillna with 'method'
is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill()
instead.
    dfx.store4.fillna(method='ffill')

```

```
[49]:
```

apple	20.0
banana	20.0
kiwi	20.0
grapes	20.0
mango	20.0
watermelon	18.0
oranges	18.0

Name: store4, dtype: float64

```
[50]: dfx.store4.fillna(method='bfill')
```

```

/tmp/ipykernel_44487/1219575873.py:1: FutureWarning: Series.fillna with 'method'
is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill()

```

```
instead.  
dfx.store4.fillna(method='bfill')
```

```
[50]: apple      20.0  
      banana     18.0  
      kiwi       18.0  
      grapes     18.0  
      mango      18.0  
      watermelon 18.0  
      oranges    NaN  
      Name: store4, dtype: float64
```

3.10 Filling with index labels

```
[51]: to_fill = pd.Series([14, 23, 12], index=['apple', 'mango', 'oranges'])  
      to_fill
```

```
[51]: apple      14  
      mango      23  
      oranges    12  
      dtype: int64
```

```
[52]: dfx.store4.fillna(to_fill)
```

```
[52]: apple      20.0  
      banana     NaN  
      kiwi       NaN  
      grapes     NaN  
      mango      23.0  
      watermelon 18.0  
      oranges    12.0  
      Name: store4, dtype: float64
```

```
[53]: dfx.fillna(dfx.mean())
```

```
[53]:
```

	store1	store2	store3	store4	store5
apple	15.0	16.0	17.0	20.0	NaN
banana	18.0	19.0	20.0	19.0	NaN
kiwi	21.0	22.0	23.0	19.0	NaN
grapes	24.0	25.0	26.0	19.0	NaN
mango	27.0	28.0	29.0	19.0	NaN
watermelon	15.0	16.0	17.0	18.0	NaN
oranges	20.0	21.0	22.0	19.0	NaN

3.11 Interpolation of missing values

```
[54]: ser3 = pd.Series([100, np.nan, np.nan, np.nan, 292])  
      ser3.interpolate()
```

```
[54]: 0    100.0  
      1    148.0  
      2    196.0  
      3    244.0  
      4    292.0  
      dtype: float64
```

```
[55]: from datetime import datetime  
      ts = pd.Series([10, np.nan, np.nan, 9],  
                     index=[datetime(2019, 1,1),  
                           datetime(2019, 2,1),  
                           datetime(2019, 3,1),  
                           datetime(2019, 5,1)])  
  
      ts
```

```
[55]: 2019-01-01    10.0  
      2019-02-01     NaN  
      2019-03-01     NaN  
      2019-05-01     9.0  
      dtype: float64
```

```
[56]: ts.interpolate()
```

```
[56]: 2019-01-01    10.000000  
      2019-02-01     9.666667  
      2019-03-01     9.333333  
      2019-05-01     9.000000  
      dtype: float64
```

```
[57]: ts.interpolate(method='time')
```

```
[57]: 2019-01-01    10.000000  
      2019-02-01     9.741667  
      2019-03-01     9.508333  
      2019-05-01     9.000000  
      dtype: float64
```

4 Discretization and binning

```
[61]: import pandas as pd
```

```
height = [120, 122, 125, 127, 121, 123, 137, 131, 161, 145, 141, 132]
```

```
bins = [118, 125, 135, 160, 200]
```

```
category = pd.cut(height, bins)
```

```
category
```

```
[61]: [(118, 125], (118, 125], (118, 125], (125, 135], (118, 125], ..., (125, 135],
(160, 200], (135, 160], (135, 160], (125, 135]]
Length: 12
Categories (4, interval[int64, right]): [(118, 125] < (125, 135] < (135, 160] <
(160, 200]]
```

```
[62]: pd.value_counts(category)
```

```
/tmp/ipykernel_44487/2654243906.py:1: FutureWarning: pandas.value_counts is
deprecated and will be removed in a future version. Use
pd.Series(obj).value_counts() instead.
```

```
pd.value_counts(category)
```

```
[62]: (118, 125]    5
(125, 135]    3
(135, 160]    3
(160, 200]    1
Name: count, dtype: int64
```

```
[63]: category2 = pd.cut(height, [118, 126, 136, 161, 200], right=False)
```

```
category2
```

```
[63]: [(118, 126), (118, 126), (118, 126), (126, 136), (118, 126), ..., (126, 136),
[161, 200), [136, 161), [136, 161), (126, 136)]
Length: 12
Categories (4, interval[int64, left]): [(118, 126) < [126, 136) < [136, 161) <
[161, 200))
```

```
[64]: bin_names = ['Short Height', 'Averge height', 'Good Height', 'Taller']
pd.cut(height, bins, labels=bin_names)
```

```
[64]: ['Short Height', 'Short Height', 'Short Height', 'Averge height', 'Short
Height', ..., 'Averge height', 'Taller', 'Good Height', 'Good Height', 'Averge
height']
```

```
Length: 12
Categories (4, object): ['Short Height' < 'Averge height' < 'Good Height' <
'Taller']
```

```
[65]: # Number of bins as integer
import numpy as np

pd.cut(np.random.rand(40), 5, precision=2)
```

```
[65]: [(0.79, 0.99], (0.2, 0.4], (0.4, 0.59], (0.59, 0.79], (-0.00095, 0.2], ...,
(0.4, 0.59], (0.4, 0.59], (-0.00095, 0.2], (0.59, 0.79], (-0.00095, 0.2]]
Length: 40
Categories (5, interval[float64, right]): [(-0.00095, 0.2] < (0.2, 0.4] < (0.4,
0.59] < (0.59, 0.79] < (0.79, 0.99]]
```

```
[66]: randomNumbers = np.random.rand(2000)
category3 = pd.qcut(randomNumbers, 4) # cut into quartiles
category3
```

```
[66]: [(0.503, 0.747], (0.252, 0.503], (0.503, 0.747], (-0.000865, 0.252], (0.747,
1.0], ..., (-0.000865, 0.252], (-0.000865, 0.252], (0.747, 1.0], (-0.000865,
0.252], (-0.000865, 0.252]]
Length: 2000
Categories (4, interval[float64, right]): [(-0.000865, 0.252] < (0.252, 0.503] <
(0.503, 0.747] < (0.747, 1.0]]
```

```
[67]: pd.value_counts(category3)
```

```
/tmp/ipykernel_44487/647498483.py:1: FutureWarning: pandas.value_counts is
deprecated and will be removed in a future version. Use
pd.Series(obj).value_counts() instead.
pd.value_counts(category3)
```

```
[67]: (-0.000865, 0.252]      500
(0.252, 0.503]             500
(0.503, 0.747]             500
(0.747, 1.0]               500
Name: count, dtype: int64
```

```
[68]: pd.qcut(randomNumbers, [0, 0.3, 0.5, 0.7, 1.0])
```

```
[68]: [(0.503, 0.699], (0.299, 0.503], (0.503, 0.699], (-0.000865, 0.299], (0.699,
1.0], ..., (-0.000865, 0.299], (-0.000865, 0.299], (0.699, 1.0], (-0.000865,
0.299], (-0.000865, 0.299]]
Length: 2000
Categories (4, interval[float64, right]): [(-0.000865, 0.299] < (0.299, 0.503] <
(0.503, 0.699] < (0.699, 1.0]]
```



```
[69]: df = pd.read_csv('https://raw.githubusercontent.com/PacktPublishing/
↳hands-on-exploratory-data-analysis-with-python/master/Chapter%204/sales.csv')
df.head(10)
```

```
[69]:
```

	Account	Company	Order	SKU	Country	Year	\
0	123456779	Kulas Inc	99985	s9-supercomputer	Aruba	1981	
1	123456784	GitHub	99986	s4-supercomputer	Brazil	2001	
2	123456782	Kulas Inc	99990	s10-supercomputer	Montserrat	1973	
3	123456783	My SQ Man	99999	s1-supercomputer	El Salvador	2015	
4	123456787	ABC Dogma	99996	s6-supercomputer	Poland	1970	
5	123456778	Super Sexy Dingo	99996	s9-supercomputer	Costa Rica	2004	
6	123456783	ABC Dogma	99981	s11-supercomputer	Spain	2006	
7	123456785	ABC Dogma	99998	s9-supercomputer	Belarus	2015	
8	123456778	Loolo INC	99997	s8-supercomputer	Mauritius	1999	
9	123456775	Kulas Inc	99997	s7-supercomputer	French Guiana	2004	

	Quantity	UnitPrice	transactionComplete
0	5148	545	False
1	3262	383	False
2	9119	407	True
3	3097	615	False
4	3356	91	True
5	2474	136	True
6	4081	195	False
7	6576	603	False
8	2460	36	False
9	1831	664	True

```
[70]: df.describe()
```

```
[70]:
```

	Account	Order	Year	Quantity	UnitPrice
count	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000
mean	1.234568e+08	99989.562900	1994.619800	4985.447300	355.866600
std	5.741156e+00	5.905551	14.432771	2868.949686	201.378478
min	1.234568e+08	99980.000000	1970.000000	0.000000	10.000000
25%	1.234568e+08	99985.000000	1982.000000	2505.750000	181.000000
50%	1.234568e+08	99990.000000	1995.000000	4994.000000	356.000000
75%	1.234568e+08	99995.000000	2007.000000	7451.500000	531.000000
max	1.234568e+08	99999.000000	2019.000000	9999.000000	700.000000

```
[71]: # Find values in order that exceeded
df['TotalPrice'] = df['UnitPrice'] * df['Quantity']
df.head(10)
```

```
[71]:
```

	Account	Company	Order	SKU	Country	Year	\
0	123456779	Kulas Inc	99985	s9-supercomputer	Aruba	1981	
1	123456784	GitHub	99986	s4-supercomputer	Brazil	2001	

2	123456782	Kulas Inc	99990	s10-supercomputer	Montserrat	1973
3	123456783	My SQ Man	99999	s1-supercomputer	El Salvador	2015
4	123456787	ABC Dogma	99996	s6-supercomputer	Poland	1970
5	123456778	Super Sexy Dingo	99996	s9-supercomputer	Costa Rica	2004
6	123456783	ABC Dogma	99981	s11-supercomputer	Spain	2006
7	123456785	ABC Dogma	99998	s9-supercomputer	Belarus	2015
8	123456778	Loolo INC	99997	s8-supercomputer	Mauritius	1999
9	123456775	Kulas Inc	99997	s7-supercomputer	French Guiana	2004

	Quantity	UnitPrice	transactionComplete	TotalPrice
0	5148	545	False	2805660
1	3262	383	False	1249346
2	9119	407	True	3711433
3	3097	615	False	1904655
4	3356	91	True	305396
5	2474	136	True	336464
6	4081	195	False	795795
7	6576	603	False	3965328
8	2460	36	False	88560
9	1831	664	True	1215784

```
[72]: # Find transaction exceeded 3000000
TotalTransaction = df["TotalPrice"]
TotalTransaction[np.abs(TotalTransaction) > 3000000]
```

```
[72]: 2      3711433
      7      3965328
      13     4758900
      15     5189372
      17     3989325
      ...
      9977    3475824
      9984    5251134
      9987    5670420
      9991    5735513
      9996    3018490
      Name: TotalPrice, Length: 2094, dtype: int64
```

```
[73]: df[np.abs(TotalTransaction) > 6741112]
```

	Account	Company	Order	SKU	Country	Year	\
818	123456781	Gen Power	99991	s1-supercomputer	Burkina Faso	1985	
1402	123456778	Will LLC	99985	s11-supercomputer	Austria	1990	
2242	123456770	Name IT	99997	s9-supercomputer	Myanmar	1979	
2876	123456772	Gen Power	99992	s10-supercomputer	Mali	2007	
3210	123456782	Loolo INC	99991	s8-supercomputer	Kuwait	2006	
3629	123456779	My SQ Man	99980	s3-supercomputer	Hong Kong	1994	

7674	123456781	Loolo INC	99989	s6-supercomputer	Sri Lanka	1994
8645	123456789	Gen Power	99996	s11-supercomputer	Suriname	2005
8684	123456785	Gen Power	99989	s2-supercomputer	Kenya	2013

	Quantity	UnitPrice	transactionComplete	TotalPrice
818	9693	696	False	6746328
1402	9844	695	True	6841580
2242	9804	692	False	6784368
2876	9935	679	False	6745865
3210	9886	692	False	6841112
3629	9694	700	False	6785800
7674	9882	691	False	6828462
8645	9742	699	False	6809658
8684	9805	694	False	6804670

5 Permuation and Random sampling

```
[74]: dat = np.arange(80).reshape(10,8)
      df = pd.DataFrame(dat)

      df
```

```
[74]:   0  1  2  3  4  5  6  7
0  0  1  2  3  4  5  6  7
1  8  9 10 11 12 13 14 15
2 16 17 18 19 20 21 22 23
3 24 25 26 27 28 29 30 31
4 32 33 34 35 36 37 38 39
5 40 41 42 43 44 45 46 47
6 48 49 50 51 52 53 54 55
7 56 57 58 59 60 61 62 63
8 64 65 66 67 68 69 70 71
9 72 73 74 75 76 77 78 79
```

```
[75]: sampler = np.random.permutation(10)
      sampler
```

```
[75]: array([2, 8, 7, 6, 5, 9, 3, 1, 4, 0])
```

```
[76]: df.take(sampler)
```

```
[76]:   0  1  2  3  4  5  6  7
2 16 17 18 19 20 21 22 23
8 64 65 66 67 68 69 70 71
7 56 57 58 59 60 61 62 63
6 48 49 50 51 52 53 54 55
```

```

5  40  41  42  43  44  45  46  47
9  72  73  74  75  76  77  78  79
3  24  25  26  27  28  29  30  31
1   8   9  10  11  12  13  14  15
4  32  33  34  35  36  37  38  39
0   0   1   2   3   4   5   6   7

```

```

[77]: # Random sample without replacement

df.take(np.random.permutation(len(df))[:3])

```

```

[77]:    0    1    2    3    4    5    6    7
5  40  41  42  43  44  45  46  47
9  72  73  74  75  76  77  78  79
3  24  25  26  27  28  29  30  31

```

```

[78]: # Random sample with replacement
sack = np.array([4, 8, -2, 7, 5])
sampler = np.random.randint(0, len(sack), size = 10)
sampler

```

```

[78]: array([3, 3, 3, 2, 1, 2, 1, 1, 0, 2])

```

```

[79]: draw = sack.take(sampler)
draw

```

```

[79]: array([ 7,  7,  7, -2,  8, -2,  8,  8,  4, -2])

```

6 Dummy variables

```

[80]: df = pd.DataFrame({'gender': ['female', 'female', 'male', 'unknown', 'male',
    ↪ 'female'], 'votes': range(6, 12, 1)})
df

```

```

[80]:   gender  votes
0  female     6
1  female     7
2   male     8
3 unknown     9
4   male    10
5  female    11

```

```

[81]: pd.get_dummies(df['gender'])

```

```

[81]:   female   male  unknown
0    True  False    False

```

1	True	False	False
2	False	True	False
3	False	False	True
4	False	True	False
5	True	False	False

```
[82]: dummies = pd.get_dummies(df['gender'], prefix='gender')
      dummies
```

```
[82]:   gender_female  gender_male  gender_unknown
0           True          False           False
1           True          False           False
2          False           True           False
3          False          False            True
4          False           True           False
5           True          False           False
```

```
[83]: with_dummy = df[['votes']].join(dummies)
      with_dummy
```

```
[83]:   votes  gender_female  gender_male  gender_unknown
0      6           True          False           False
1      7           True          False           False
2      8          False           True           False
3      9          False          False            True
4     10          False           True           False
5     11           True          False           False
```

6.1 Benefits of data transformation

1. Data transformation promotes interoperability between several applications. The main reason for creating a similar format and structure in the dataset is that it becomes compatible with other systems.
2. Comprehensibility for both humans and computers is improved when using better-organized data compared to messier data.
3. Data transformation ensures a higher degree of data quality and protects applications from several computational challenges such as null values, unexpected duplicates, and incorrect indexings, as well as incompatible structures or formats.
4. Data transformation ensures higher performance and scalability for modern analytical databases and dataframes.