

# EDA\_with\_Personal\_Email\_Analysis

November 27, 2024

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
[2]: !pip install mailbox
```

```
Collecting mailbox
  Downloading mailbox-0.4.tar.gz (4.1 kB)
  Preparing metadata (setup.py) ... done
Building wheels for collected packages: mailbox
  Building wheel for mailbox (setup.py) ... done
  Created wheel for mailbox: filename=mailbox-0.4-py3-none-any.whl
size=4684
sha256=8fac459ede30179c2914f0609ad51234c7551ca8505d9900bb8fadc42302e1f9
  Stored in directory: /home/test/.cache/pip/wheels/06/cd/9a/64b75da2511d797260d
3b3cb8cfbf66e700119cc045a9be2c9
Successfully built mailbox
Installing collected packages: mailbox
Successfully installed mailbox-0.4
```

```
[ ]: from google.colab import drive
drive.mount('/content/gdrive')
```

```
[ ]: import mailbox
mboxfile = "gdrive/My Drive/Colab Notebooks/gmail.mbox"

mbox = mailbox.mbox(mboxfile)
mbox
```

```
[ ]: for key in mbox[0].keys():
    print(key)
```

```
[ ]: import csv

with open('mailbox.csv', 'w') as outputfile:
    writer = csv.writer(outputfile)
    writer.writerow(['subject', 'from', 'date', 'to', 'label', 'thread'])
```

```
for message in mbox:
    writer.writerow([message['subject'], message['from'], message['date'],
↪message['to'], message['X-Gmail-Labels'], message['X-GM-THRID']])
```

```
[ ]: dfs = pd.read_csv('mailbox.csv', names=['subject', 'from', 'date', 'to',
↪'label', 'thread'])
```

```
[ ]: dfs.dtypes
```

```
[ ]: dfs['date'] = dfs['date'].apply(lambda x: pd.to_datetime(x, errors='coerce',
↪utc=True))
```

```
[ ]: dfs = dfs[dfs['date'].notna()]
```

```
[ ]: dfs.to_csv('gmail.csv')
```

```
[ ]: dfs.info()
```

```
[ ]: dfs.head(10)
```

```
[ ]: dfs.columns
```

```
[ ]: import re

def extract_email_ID(string):
    email = re.findall(r'<(.*?)>', string)
    if not email:
        email = list(filter(lambda y: '@' in y, string.split()))
    return email[0] if email else np.nan
```

```
[ ]: dfs['from'] = dfs['from'].apply(lambda x: extract_email_ID(x))
```

```
[ ]: myemail = 'itsmeskm99@gmail.com'
dfs['label'] = dfs['from'].apply(lambda x: 'sent' if x==myemail else 'inbox')
```

```
[ ]: dfs.drop(columns='to', inplace=True)
```

```
[ ]: dfs.head(10)
```

```
[ ]: import datetime
import pytz

def refactor_timezone(x):
    est = pytz.timezone('US/Eastern')
    return x.astimezone(est)
```

```
[ ]: dfs['date'] = dfs['date'].apply(lambda x: refactor_timezone(x))
```

```
[ ]: dfs['dayofweek'] = dfs['date'].apply(lambda x: x.weekday_name)
dfs['dayofweek'] = pd.Categorical(dfs['dayofweek'], categories=[
    'Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',
    'Saturday', 'Sunday'], ordered=True)

[ ]: dfs['timeofday'] = dfs['date'].apply(lambda x: x.hour + x.minute/60 + x.second/
    ↪3600)

[ ]: dfs['hour'] = dfs['date'].apply(lambda x: x.hour)

[ ]: dfs['year_int'] = dfs['date'].apply(lambda x: x.year)

[ ]: dfs['year'] = dfs['date'].apply(lambda x: x.year + x.dayofyear/365.25)

[ ]: dfs.index = dfs['date']
del dfs['date']

[ ]: print(dfs.index.min().strftime('%a, %d %b %Y %I:%M %p'))
print(dfs.index.max().strftime('%a, %d %b %Y %I:%M %p'))

print(dfs['label'].value_counts())

[ ]: import matplotlib.pyplot as plt
from matplotlib.ticker import MaxNLocator

[ ]: def plot_todo_vs_year(df, ax, color='C0', s=0.5, title=''):
    ind = np.zeros(len(df), dtype='bool')
    est = pytz.timezone('US/Eastern')

    df[~ind].plot.scatter('year', 'timeofday', s=s, alpha=0.6, ax=ax, color=color)
    ax.set_ylim(0, 24)
    ax.yaxis.set_major_locator(MaxNLocator(8))
    ax.set_yticklabels([datetime.datetime.strptime(str(int(np.mod(ts, 24))),
    ↪"%H").strftime("%I %p") for ts in ax.get_yticks()]);

    ax.set_xlabel('')
    ax.set_ylabel('')
    ax.set_title(title)
    ax.grid(ls=':', color='k')

    return ax

[ ]: sent = dfs[dfs['label']=='sent']
received = dfs[dfs['label']=='inbox']
```

```
[ ]: fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(15, 4))
```

```
plot_todo_vs_year(sent, ax[0], title='Sent')
plot_todo_vs_year(received, ax[1], title='Received')
```

```
[ ]: def plot_number_perday_per_year(df, ax, label=None, dt=0.3, **plot_kwargs):
    year = df[df['year'].notna()]['year'].values
    T = year.max() - year.min()
    bins = int(T / dt)
    weights = 1 / (np.ones_like(year) * dt * 365.25)
    ax.hist(year, bins=bins, weights=weights, label=label, **plot_kwargs);
    ax.grid(ls=':', color='k')
```

```
[ ]: from scipy import ndimage
```

```
def plot_number_perdhour_per_year(df, ax, label=None, dt=1, smooth=False,
    weight_fun=None, **plot_kwargs):

    tod = df[df['timeofday'].notna()]['timeofday'].values
    year = df[df['year'].notna()]['year'].values
    Ty = year.max() - year.min()
    T = tod.max() - tod.min()
    bins = int(T / dt)
    if weight_fun is None:
        weights = 1 / (np.ones_like(tod) * Ty * 365.25 / dt)
    else:
        weights = weight_fun(df)
    if smooth:
        hst, xedges = np.histogram(tod, bins=bins, weights=weights);
        x = np.delete(xedges, -1) + 0.5*(xedges[1] - xedges[0])
        hst = ndimage.gaussian_filter(hst, sigma=0.75)
        f = interp1d(x, hst, kind='cubic')
        x = np.linspace(x.min(), x.max(), 10000)
        hst = f(x)
        ax.plot(x, hst, label=label, **plot_kwargs)
    else:
        ax.hist(tod, bins=bins, weights=weights, label=label, **plot_kwargs);

    ax.grid(ls=':', color='k')
    orientation = plot_kwargs.get('orientation')
    if orientation is None or orientation == 'vertical':
        ax.set_xlim(0, 24)
        ax.xaxis.set_major_locator(MaxNLocator(8))
        ax.set_xticklabels([datetime.datetime.strptime(str(int(np.mod(ts, 24))), "%H").strftime("%I %p")
            for ts in ax.get_xticks()]));
```

```

elif orientation == 'horizontal':
    ax.set_ylim(0, 24)
    ax.yaxis.set_major_locator(MaxNLocator(8))
    ax.set_yticklabels([datetime.datetime.strptime(str(int(np.mod(ts,
↪24))), "%H").strftime("%I %p")
                        for ts in ax.get_yticks()]));

```

```

[ ]: class TriplePlot:
    def __init__(self):
        gs = gridspec.GridSpec(6, 6)
        self.ax1 = plt.subplot(gs[2:6, :4])
        self.ax2 = plt.subplot(gs[2:6, 4:6], sharey=self.ax1)
        plt.setp(self.ax2.get_yticklabels(), visible=False);
        self.ax3 = plt.subplot(gs[:2, :4])
        plt.setp(self.ax3.get_xticklabels(), visible=False);

    def plot(self, df, color='darkblue', alpha=0.8, markersize=0.5, yr_bin=0.1,
↪hr_bin=0.5):
        plot_todo_vs_year(df, self.ax1, color=color, s=markersize)
        plot_number_perdhour_per_year(df, self.ax2, dt=hr_bin, color=color,
↪alpha=alpha, orientation='horizontal')
        self.ax2.set_xlabel('Average emails per hour')
        plot_number_perday_per_year(df, self.ax3, dt=yr_bin, color=color,
↪alpha=alpha)
        self.ax3.set_ylabel('Average emails per day')

```

```

[ ]: import matplotlib.gridspec as gridspec
import matplotlib.patches as mpatches

plt.figure(figsize=(12,12));
tpl = TriplePlot()

tpl.plot(received, color='C0', alpha=0.5)
tpl.plot(sent, color='C1', alpha=0.5)
p1 = mpatches.Patch(color='C0', label='Incoming', alpha=0.5)
p2 = mpatches.Patch(color='C1', label='Outgoing', alpha=0.5)
plt.legend(handles=[p1, p2], bbox_to_anchor=[1.45, 0.7], fontsize=14,
↪shadow=True);

```

```

[ ]: counts = dfs.dayofweek.value_counts(sort=False)
counts.plot(kind='bar')

```

```

[ ]: addrs = received['from'].value_counts()

addrs[0:4]

```

```
[ ]: plt.figure(figsize=(12,12));

tpl = TriplePlot()

labels = []
colors = ['C{}'.format(ii) for ii in range(9)]
idx = np.array([1,2,3,7])
for ct, addr in enumerate(addr.index[idx]):
    tpl.plot(dfs[dfs['from'] == addr], color=colors[ct], alpha=0.3, yr_bin=0.5,
    markersize=1.0)
    labels.append(mpatches.Patch(color=colors[ct], label=addr[0:4], alpha=0.5))
plt.legend(handles=labels, bbox_to_anchor=[1.4, 0.9], fontsize=12, shadow=True);
```

```
[ ]: sdw = sent.groupby('dayofweek').size() / len(sent)
rdw = received.groupby('dayofweek').size() / len(received)

df_tmp = pd.DataFrame(data={'Outgoing Email': sdw, 'Incoming Email':rdw})
df_tmp.plot(kind='bar', rot=45, figsize=(8,5), alpha=0.5)
plt.xlabel('');
plt.ylabel('Fraction of weekly emails');
plt.grid(ls=':', color='k', alpha=0.5)
```

```
[ ]: import scipy.ndimage
from scipy.interpolate import interp1d

plt.figure(figsize=(8,5))
ax = plt.subplot(111)
for ct, dow in enumerate(dfs.dayofweek.cat.categories):
    df_r = received[received['dayofweek']==dow]
    weights = np.ones(len(df_r)) / len(received)
    wfun = lambda x: weights
    plot_number_perdhour_per_year(df_r, ax, dt=1, smooth=True, color=f'C{ct}',
    alpha=0.8, lw=3, label=dow, weight_fun=wfun)

    df_s = sent[sent['dayofweek']==dow]
    weights = np.ones(len(df_s)) / len(sent)
    wfun = lambda x: weights
    plot_number_perdhour_per_year(df_s, ax, dt=1, smooth=True, color=f'C{ct}',
    alpha=0.8, lw=2, label=dow, ls='--', weight_fun=wfun)
ax.set_ylabel('Fraction of weekly emails per hour')
plt.legend(loc='upper left')
```

```
[ ]: from wordcloud import WordCloud

df_no_arxiv = dfs[dfs['from'] != 'no-reply@arXiv.org']
text = ' '.join(map(str, sent['subject'].values))
```

```
[ ]: stopwords = ['Re', 'Fwd', '3A_']
wrld = WordCloud(width=700, height=480, margin=0, collocations=False)
for sw in stopwords:
    wrld.stopwords.add(sw)
wordcloud = wrld.generate(text)

plt.figure(figsize=(25,15))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis("off")
plt.margins(x=0, y=0)
```