

Exploratory Data Analysis Wine Quality dataset

November 27, 2024

1 Exploratory Data Analysis Wine Quality dataset

Agenda

- Loading the dataset
- Data wrangling for missing variables.
- Data transformation.
- Data visualization
- Answering the main questions:

```
[ ]:
```

```
[ ]: import pandas as pd
```

```
[ ]: df_red = pd.read_csv("https://archive.ics.uci.edu/ml/machine-learning-databases/  
    ↪ wine-quality/winequality-red.csv", delimiter=";")  
df_white = pd.read_csv("https://archive.ics.uci.edu/ml/  
    ↪ machine-learning-databases/wine-quality/winequality-white.csv", delimiter=";  
    ↪")
```

```
[115]: df_red.columns
```

```
[115]: Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',  
          'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',  
          'pH', 'sulphates', 'alcohol', 'quality'],  
          dtype='object')
```

```
[116]: df_white.columns
```

```
[116]: Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',  
          'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',  
          'pH', 'sulphates', 'alcohol', 'quality'],  
          dtype='object')
```

```
[117]: df_red.dtypes
```

```
[117]: fixed acidity          float64  
       volatile acidity    float64
```

```

citric acid          float64
residual sugar       float64
chlorides            float64
free sulfur dioxide  float64
total sulfur dioxide float64
density             float64
pH                  float64
sulphates           float64
alcohol             float64
quality             int64
dtype: object

```

```
[118]: df_red.iloc[100:110]
```

```

[118]:      fixed acidity  volatile acidity  citric acid  ...  sulphates  alcohol
quality
100          8.3          0.610          0.30  ...      0.61      10.2
6
101          7.8          0.500          0.30  ...      0.56      10.4
6
102          8.1          0.545          0.18  ...      0.59       9.0
6
103          8.1          0.575          0.22  ...      0.51       9.2
5
104          7.2          0.490          0.24  ...      0.48       9.4
5
105          8.1          0.575          0.22  ...      0.51       9.2
5
106          7.8          0.410          0.68  ...      1.31       9.3
5
107          6.2          0.630          0.31  ...      0.79       9.3
5
108          8.0          0.330          0.53  ...      0.80       9.6
6
109          8.1          0.785          0.52  ...      0.69       9.3
5

```

```
[10 rows x 12 columns]
```

```
[119]: df_red.describe()
```

```

[119]:      fixed acidity  volatile acidity  ...  alcohol  quality
count    1599.000000    1599.000000  ...  1599.000000  1599.000000
mean       8.319637       0.527821  ...    10.422983    5.636023
std       1.741096       0.179060  ...     1.065668    0.807569
min       4.600000       0.120000  ...     8.400000    3.000000
25%       7.100000       0.390000  ...     9.500000    5.000000

```

50%	7.900000	0.520000	...	10.200000	6.000000
75%	9.200000	0.640000	...	11.100000	6.000000
max	15.900000	1.580000	...	14.900000	8.000000

[8 rows x 12 columns]

```
[120]: df_red.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
fixed acidity          1599 non-null float64
volatile acidity       1599 non-null float64
citric acid            1599 non-null float64
residual sugar         1599 non-null float64
chlorides              1599 non-null float64
free sulfur dioxide    1599 non-null float64
total sulfur dioxide   1599 non-null float64
density                1599 non-null float64
pH                     1599 non-null float64
sulphates              1599 non-null float64
alcohol                1599 non-null float64
quality                1599 non-null int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

2 Analysis of Red Wine

```
[121]: import seaborn as sns
```


```
sns.set(rc={'figure.figsize': (14, 8)})
sns.countplot(df_red['quality'])
```

```
[121]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc7e56e31d0>
```




```
[122]: sns.pairplot(df_red)
```

```
[122]: <seaborn.axisgrid.PairGrid at 0x7fc7e56a0400>
```



```
[123]: sns.heatmap(df_red.corr(), annot=True, fmt='.2f', linewidths=2)
```


```
[123]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc7e2f6b470>
```



```
[ ]:
```

```
[124]: sns.distplot(df_red['alcohol'])
```

```
[124]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc7e2549668>
```




```
[125]: from scipy.stats import skew  
  
skew(df_red['alcohol'])
```

```
[125]: 0.8600210646566755
```


```
[126]: sns.boxplot(x='quality', y='alcohol', data = df_red)
```

```
[126]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc7e2417940>
```




```
[127]: sns.boxplot(x='quality', y='alcohol', data = df_red, showfliers=False)
```

```
[127]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc7e23be748>
```



```
[128]: sns.jointplot(x='alcohol',y='pH',data=df_red, kind='reg')
```


[128]: <seaborn.axisgrid.JointGrid at 0x7fc7e246e470>



```
[ ]: from scipy.stats import pearsonr

def get_correlation(column1, column2, df):
    pearson_corr, p_value = pearsonr(df[column1], df[column2])
    print("Correlation between {} and {} is {}".format(column1, column2,
    ↪pearson_corr))
    print("P-value of this correlation is {}".format(p_value))
```

```
[130]: get_correlation('alcohol', 'pH', df_red)
```

Correlation between alcohol and pH is 0.20563250850549825

P-value of this correlation is 9.96449774146556e-17

```
[131]: df_white.describe()
```

```
[131]:
```

	fixed acidity	volatile acidity	...	alcohol	quality
count	4898.000000	4898.000000	...	4898.000000	4898.000000
mean	6.854788	0.278241	...	10.514267	5.877909
std	0.843868	0.100795	...	1.230621	0.885639
min	3.800000	0.080000	...	8.000000	3.000000
25%	6.300000	0.210000	...	9.500000	5.000000
50%	6.800000	0.260000	...	10.400000	6.000000
75%	7.300000	0.320000	...	11.400000	6.000000
max	14.200000	1.100000	...	14.200000	9.000000

[8 rows x 12 columns]

3 White wine analysis

```
[153]: print("white mean = ",df_white["quality"].mean())  
print("red mean =",df_red["quality"].mean())
```

white mean = 5.87790935075541

red mean = 5.6360225140712945

```
[154]: d = {'color': ['red','white'], 'mean_quality': [5.636023,5.877909]}  
df_mean = pd.DataFrame(data=d)  
df_mean
```

```
[154]:
```

	color	mean_quality
0	red	5.636023
1	white	5.877909

```
[ ]: # Let us add new attribute called wine_category to the both dataframe
```

```
df_white['wine_category'] = 'white'  
df_red['wine_category'] = 'red'
```

```
[158]: print('RED WINE: List of "quality"', sorted(df_red['quality'].unique()))  
print('WHITE WINE: List of "quality"', sorted(df_white['quality'].unique()))
```

RED WINE: List of "quality" [3, 4, 5, 6, 7, 8]

WHITE WINE: List of "quality" [3, 4, 5, 6, 7, 8, 9]

4 Convert into categorical dataset

```
[ ]: df_red['quality_label'] = df_red['quality'].apply(lambda value: ('low' if value
    ↪ <= 5 else 'medium') if value <= 7 else 'high')
df_red['quality_label'] = pd.Categorical(df_red['quality_label'],
    ↪ categories=['low', 'medium', 'high'])

df_white['quality_label'] = df_white['quality'].apply(lambda value: ('low' if
    ↪ value <= 5 else 'medium') if value <= 7 else 'high')
df_white['quality_label'] = pd.Categorical(df_white['quality_label'],
    ↪ categories=['low', 'medium', 'high'])
```

```
[163]: print(df_white['quality_label'].value_counts())
df_red['quality_label'].value_counts()
```

```
medium    3078
low        1640
high       180
Name: quality_label, dtype: int64
```

```
[163]: medium    837
low        744
high       18
Name: quality_label, dtype: int64
```

```
[170]: df_wines = pd.concat([df_red, df_white])

# Re-shuffle records just to randomize data points.
# `drop=True`: this resets the index to the default integer index.
df_wines = df_wines.sample(frac=1.0, random_state=42).reset_index(drop=True)
df_wines.head(10)
```

```
[170]:   fixed acidity  volatile acidity  ...  wine_category  quality_label
0           7.0           0.17  ...         white         high
1           7.7           0.64  ...          red         low
2           6.8           0.39  ...         white        medium
3           6.3           0.28  ...         white        medium
4           7.4           0.35  ...         white        medium
5           7.2           0.53  ...          red        medium
6           7.5           0.27  ...         white         low
7           6.8           0.11  ...         white        medium
8           9.0           0.44  ...          red         low
9           7.1           0.23  ...         white        medium
```

```
[10 rows x 14 columns]
```

```
[168]: subset_attr = ['alcohol', 'density', 'pH', 'quality']

low = round(df_wines[df_wines['quality_label'] == 'low'][subset_attr].
↳describe(), 2)
medium = round(df_wines[df_wines['quality_label'] == 'medium'][subset_attr].
↳describe(), 2)
high = round(df_wines[df_wines['quality_label'] == 'high'][subset_attr].
↳describe(), 2)

pd.concat([low, medium, high], axis=1,
          keys=[' Low Quality Wine',
                ' Medium Quality Wine',
                ' High Quality Wine'])
```

```
[168]:
```


	Low Quality Wine			...	High Quality Wine	
	alcohol	density	pH	...	density	pH
quality						
count	2384.00	2384.00	2384.00	...	198.00	198.00
198.00						
mean	9.87	1.00	3.21	...	0.99	3.23
8.03						
std	0.84	0.00	0.16	...	0.00	0.16
0.16						
min	8.00	0.99	2.74	...	0.99	2.88
8.00						
25%	9.30	0.99	3.11	...	0.99	3.13
8.00						
50%	9.60	1.00	3.20	...	0.99	3.23
8.00						
75%	10.40	1.00	3.31	...	0.99	3.33
8.00						
max	14.90	1.00	3.90	...	1.00	3.72
9.00						

[8 rows x 12 columns]

```
[ ]: import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
%matplotlib inline
```

```
[176]: fig = df_wines.hist(bins=15, color='fuchsia', edgecolor='darkmagenta',
↳linewidth=1.0, xlabelsize=10, ylabelsize=10, xrot=45, yrot=0,
↳figsize=(10,9), grid=False)

plt.tight_layout(rect=(0, 0, 1.5, 1.5))
```



```
[160]: df_red.head()
```

```
[160]:   fixed acidity  volatile acidity  ...  wine_category  quality_label
0          7.4           0.70  ...          red          low
1          7.8           0.88  ...          red          low
2          7.8           0.76  ...          red          low
3         11.2           0.28  ...          red        medium
4          7.4           0.70  ...          red          low
```

[5 rows x 14 columns]

```
[185]: fig, (ax) = plt.subplots(1, 1, figsize=(14,8))

hm = sns.heatmap(df_wines.corr(),
                  ax=ax,
                  cmap="bwr",
```


```

        annot=True,
        fmt='.2f',
        linewidths=.05)

fig.subplots_adjust(top=0.94)
fig.suptitle('Combined Wine Attributes and their Correlation Heatmap',
            fontsize=14,
            fontweight='bold')

```

[185]: `Text(0.5, 0.98, 'Combined Wine Attributes and their Correlation Heatmap')`



4.1 Discrete categorical attributes


```

[188]: fig = plt.figure(figsize=(16, 8))

sns.countplot(data=df_wines, x="quality", hue="wine_category")

```

[188]: `<matplotlib.axes._subplots.AxesSubplot at 0x7fc7e196da20>`




4.2 3D visualization

```
[196]: fig = plt.figure(figsize=(16, 12))
ax = fig.add_subplot(111, projection='3d')

xscale = df_wines['residual sugar']
yscale = df_wines['free sulfur dioxide']
zscale = df_wines['total sulfur dioxide']
ax.scatter(xscale, yscale, zscale, s=50, alpha=0.6, edgecolors='w')

ax.set_xlabel('Residual Sugar')
ax.set_ylabel('free sulfur dioxide')
ax.set_zlabel('Total sulfur dioxide')

plt.show()
```




```
[200]: fig = plt.figure(figsize=(16, 12))

plt.scatter(x = df_wines['fixed acidity'],
            y = df_wines['free sulfur dioxide'],
            s = df_wines['total sulfur dioxide'] * 2,
            alpha=0.4,
            edgecolors='w')

plt.xlabel('Fixed Acidity')
plt.ylabel('free sulfur dioxide')
plt.title('Wine free sulfur dioxide Content - Fixed Acidity - total sulfur_
↪dioxide', y=1.05)
```

```
[200]: Text(0.5, 1.05, 'Wine free sulfur dioxide Content - Fixed Acidity - total sulfur
dioxide')
```

```
[195]: df_wines.columns
```

```
[195]: Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
          'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
          'pH', 'sulphates', 'alcohol', 'quality', 'wine_category',
          'quality_label'],
          dtype='object')
```

```
[ ]: from sklearn.linear_model import LogisticRegression
      from sklearn.svm import LinearSVC,SVC
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.ensemble import
        ↳ RandomForestClassifier, GradientBoostingClassifier, AdaBoostClassifier
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.naive_bayes import GaussianNB
```

```
[ ]: from sklearn.model_selection import train_test_split, cross_validate
      from sklearn.preprocessing import MinMaxScaler, StandardScaler, LabelEncoder
      from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
[ ]: label_quality = LabelEncoder()

df_wines['quality_label'] = label_quality.
↳fit_transform(df_wines['quality_label'])
```

```
[236]: df_wines.tail(10)
```

```
[236]:      fixed acidity  volatile acidity  ...  wine_category  quality_label
6487           6.1             0.22  ...             1             2
6488          10.3             0.50  ...             0             2
6489           6.4             0.31  ...             1             1
6490           5.9             0.26  ...             1             1
6491           8.0             0.34  ...             1             0
6492           7.6             0.32  ...             1             1
6493           5.6             0.28  ...             1             2
6494           6.4             0.37  ...             1             1
6495           6.5             0.26  ...             1             1
6496           7.2             0.62  ...             0             1
```

[10 rows x 14 columns]

```
[ ]: x_train,x_test,y_train,y_test=train_test_split(df_wines.
↳drop(['quality', 'wine_category'],axis=1),df_wines['quality_label'],test_size=0.
↳30,random_state=42)

models=[LogisticRegression(),
        LinearSVC(),
        SVC(kernel='rbf'),
        KNeighborsClassifier(),
        RandomForestClassifier(),
        DecisionTreeClassifier(),
        GradientBoostingClassifier(),
        GaussianNB()]
```

```
[242]: model_names=['LogisticRegression',
                    'LinearSVM',
                    'rbfSVM',
                    'KNearestNeighbors',
                    'RandomForestClassifier',
                    'DecisionTree',
                    'GradientBoostingClassifier',
                    'GaussianNB']

acc=[]
eval_acc={}

for model in range(len(models)):
```

```

classification_model=models[model]
classification_model.fit(x_train,y_train)
pred=classification_model.predict(x_test)
acc.append(accuracy_score(pred,y_test))

eval_acc={'Modelling Algorithm':model_names,'Accuracy':acc}
eval_acc

```

```

/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/_logistic.py:940:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
/usr/local/lib/python3.6/dist-packages/sklearn/svm/_base.py:947:
ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
"the number of iterations.", ConvergenceWarning)

```

```

[242]: {'Accuracy': [0.9687179487179487,
0.9733333333333334,
0.6051282051282051,
0.6912820512820513,
1.0,
1.0,
1.0,
1.0],
'Modelling Algorithm': ['LogisticRegression',
'LinearSVM',
'rbfSVM',
'KNearestNeighbors',
'RandomForestClassifier',
'DecisionTree',
'GradientBoostingClassifier',
'GaussianNB']}

```

```

[239]: acc_table=pd.DataFrame(eval_acc)
acc_table = acc_table.sort_values(by='Accuracy', ascending=[False])
acc_table

```

```

[239]:
Modelling Algorithm  Accuracy
4  RandomForestClassifier  1.000000
5           DecisionTree  1.000000

```

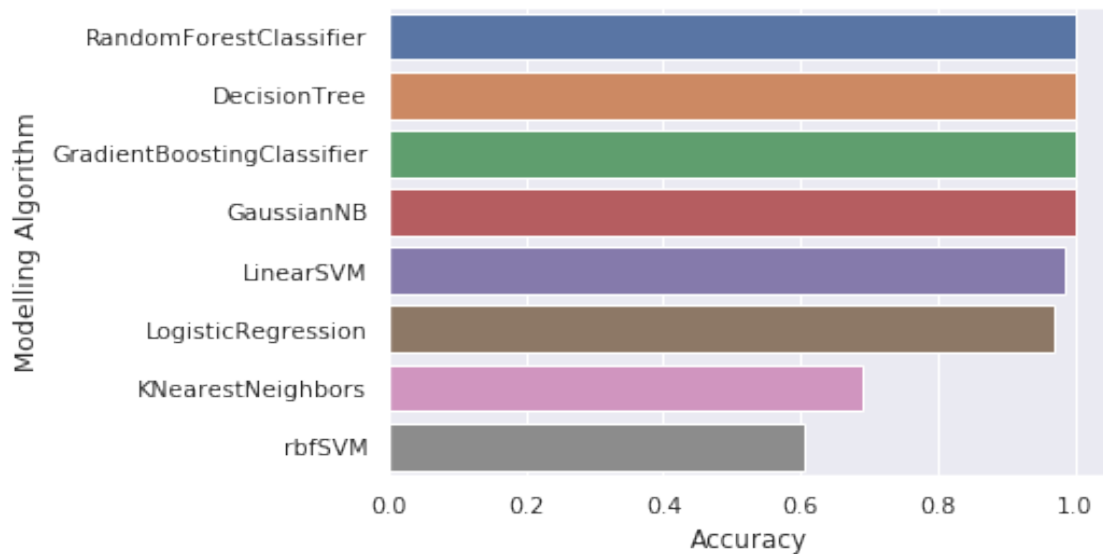
```

6 GradientBoostingClassifier 1.000000
7 GaussianNB 1.000000
1 LinearSVM 0.983590
0 LogisticRegression 0.968718
3 KNearestNeighbors 0.691282
2 rbfSVM 0.605128

```

```
[240]: sns.barplot(y='Modelling Algorithm',x='Accuracy',data=acc_table)
```

```
[240]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc7e0624c50>
```



```
[241]: sns.catplot(x='Modelling_
↳Algorithm',y='Accuracy',data=acc_table,kind='point',size=4,aspect=3.5)
```

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:3695: UserWarning:
The `size` parameter has been renamed to `height`; please update your code.
warnings.warn(msg, UserWarning)

```

```
[241]: <seaborn.axisgrid.FacetGrid at 0x7fc7e0efab38>
```

